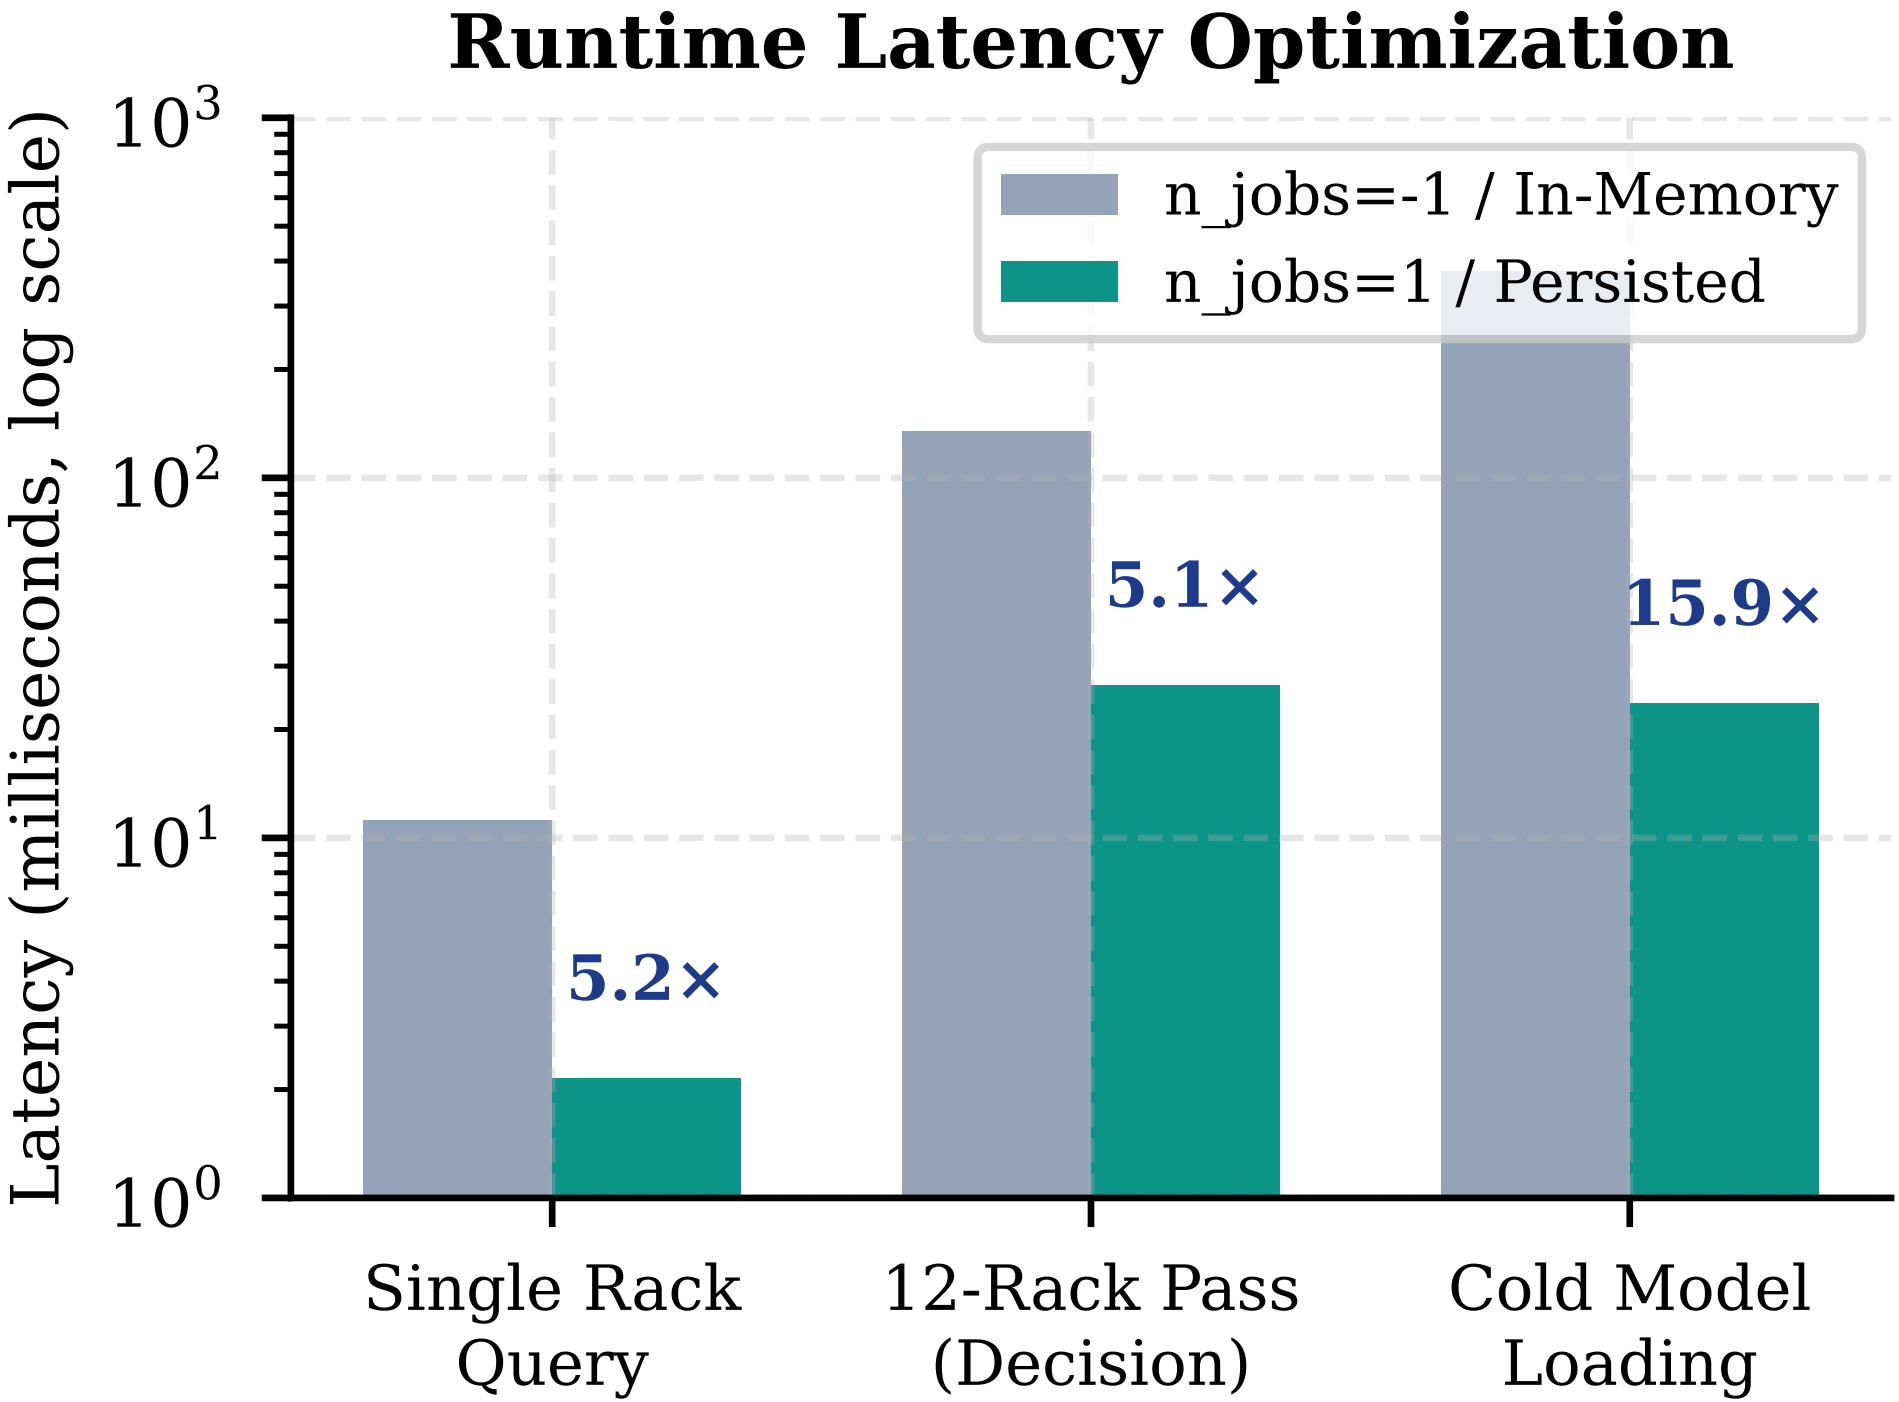
Runtime Latency Optimization
Latency (milliseconds, log scale)
10^3
10^2
10^1
10^0
n_jobs=-1 / In-Memory
n_jobs=1 / Persisted
5.2×
5.1×
15.9×
Single Rack Query
12-Rack Pass (Decision)
Cold Model Loading